

# Unsupervised Learning

## K-MEANS

Dr. Amani RAAD

[amani.raad@ul.edu.lb](mailto:amani.raad@ul.edu.lb)

2022-2023

# Supervised learning vs. unsupervised learning


- **Supervised learning:** discover patterns in the data that relate data attributes with a target (class) attribute.
  - These patterns are then utilized to predict the values of the target attribute in future data instances.
- **Unsupervised learning:** The data have no target attribute.
  - We want to explore the data to find some intrinsic structures in them.

# Clustering

- Clustering is a technique for finding **similarity groups** in data, called **clusters**. I.e.,
  - it groups data instances that are similar to (near) each other in one cluster and data instances that are very different (far away) from each other into different clusters.
- Clustering is often called an **unsupervised learning** task as no class values denoting an *a priori* grouping of the data instances are given, which is the case in supervised learning.
- Due to historical reasons, clustering is often considered synonymous with unsupervised learning.
  - In fact, association rule mining is also unsupervised
- This chapter focuses on clustering.

# An illustration

- The data set has three natural groups of data points, i.e., 3 natural clusters.



# What is clustering for?

- Let us see some real-life examples
- **Example 1:** groups people of similar sizes together to make “small”, “medium” and “large” T-Shirts.
  - Tailor-made for each person: too expensive
  - One-size-fits-all: does not fit all.
- **Example 2:** In marketing, segment customers according to their similarities
  - To do targeted marketing.

# What is clustering for? (cont...)

- **Example 3:** Given a collection of text documents, we want to organize them according to their content similarities,
  - To produce a topic hierarchy
- **In fact, clustering is one of the most utilized data mining techniques.**
  - It has a long history, and used in almost every field, e.g., medicine, psychology, botany, sociology, biology, archeology, marketing, insurance, libraries, etc.
  - In recent years, due to the rapid increase of online documents, text clustering becomes important.

# Aspects of clustering

- A clustering algorithm
  - Partitional clustering
  - Hierarchical clustering
  - ...
- A distance (similarity, or dissimilarity) function
- Clustering quality
  - Inter-clusters distance  $\Rightarrow$  maximized
  - Intra-clusters distance  $\Rightarrow$  minimized
- The **quality** of a clustering result depends on the algorithm, the distance function, and the application.

# K-means clustering

- K-means is a **partitional clustering** algorithm
- Let the set of data points (or instances)  $D$  be

$$\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\},$$

where  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ir})$  is a **vector** in a real-valued space  $X \subseteq R^r$ , and  $r$  is the number of attributes (dimensions) in the data.

- The  $k$ -means algorithm partitions the given data into  $k$  clusters.
  - Each cluster has a cluster **center**, called **centroid**.
  - $k$  is specified by the user



# K-means algorithm

- Given  $k$ , the *k-means* algorithm works as follows:
  - 1) Randomly choose  $k$  data points (**seeds**) to be the initial **centroids**, cluster centers
  - 2) Assign each data point to the closest **centroid**
  - 3) Re-compute the **centroids** using the current cluster memberships.
  - 4) If a convergence criterion is not met, go to **2**).

# K-means algorithm – (cont ...)

**Algorithm**  $k$ -means( $k, D$ )

- 1 Choose  $k$  data points as the initial centroids (cluster centers)
- 2 **repeat**
- 3     **for** each data point  $\mathbf{x} \in D$  **do**
- 4         compute the distance from  $\mathbf{x}$  to each centroid;
- 5         assign  $\mathbf{x}$  to the closest centroid         // a centroid represents a cluster
- 6     **endfor**
- 7     re-compute the centroids using the current cluster memberships
- 8 **until** the stopping criterion is met


# Stopping/convergence criterion

1. no (or minimum) re-assignments of data points to different clusters,
2. no (or minimum) change of centroids, or
3. minimum decrease in the **sum of squared error** (SSE) or WCSS,


$$WCSS = \sum_{j=1}^k \sum_{\mathbf{x} \in C_j} \text{dist}(\mathbf{x}, \mathbf{m}_j)^2 \quad (1)$$

- $C_j$  is the  $j$ th cluster,  $\mathbf{m}_j$  is the centroid of cluster  $C_j$  (the mean vector of all the data points in  $C_j$ ), and  $\text{dist}(\mathbf{x}, \mathbf{m}_j)$  is the distance between data point  $\mathbf{x}$  and centroid  $\mathbf{m}_j$ .


# An example



(A). Random selection of  $k$  centers




Iteration 1: (B). Cluster assignment




(C). Re-compute centroids


## An example (cont ...)




*Iteration 2: (D). Cluster assignment*



*(E). Re-compute centroids*



*Iteration 3: (F). Cluster assignment*



*(G). Re-compute centroids*

# An example distance function


The  $k$ -means algorithm can be used for any application data set where the **mean** can be defined and computed. In the **Euclidean space**, the mean of a cluster is computed with:

$$\mathbf{m}_j = \frac{1}{|C_j|} \sum_{\mathbf{x}_i \in C_j} \mathbf{x}_i \quad (2)$$

where  $|C_j|$  is the number of data points in cluster  $C_j$ . The distance from one data point  $\mathbf{x}_i$  to a mean (centroid)  $\mathbf{m}_j$  is computed with

$$\begin{aligned} dist(\mathbf{x}_i, \mathbf{m}_j) &= \| \mathbf{x}_i - \mathbf{m}_j \| \\ &= \sqrt{(x_{i1} - m_{j1})^2 + (x_{i2} - m_{j2})^2 + \dots + (x_{ir} - m_{jr})^2} \end{aligned} \quad (3)$$

# Other geometrical distances



$$d_E(\vec{p}, \vec{q}) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

Euclidean distance

$$d_M(\vec{p}, \vec{q}) = \sum_{i=1}^n |p_i - q_i|$$

Manhattan distance

$$d_{Cheb}(\vec{p}, \vec{q}) = \max_i |(p_i - q_i)|$$

Maximum(Chebychev) distance



# Other distances: Mahalanobis

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad c = \begin{bmatrix} c_1 \\ c_2 \end{bmatrix}$$

$\sigma_1^2$  - Horizontal variance

$\sigma_2^2$  - Vertical variance

$$d(x, c) = \sqrt{\frac{(x_1 - c_1)^2}{\sigma_1^2} + \frac{(x_2 - c_2)^2}{\sigma_2^2}}$$

Mahalanobis distance for a two dimensional vector with no covariance

# Remarks:

- When you use Euclidean distance, you assume that the clusters have identity covariances. In 2D, this means that your clusters have circular shapes. Obviously, if the covariances of the natural groupings in your data are not identity matrices, e.g. in 2D, clusters have elliptical shaped covariances, then using Mahalanobis over Euclidean will be much better modeling.
- Manhattan distance should give more robust results, whereas Euclidean distance is likely to be influenced by outliers.

# Remarks:

- **The value of distance measures is intimately related to the scale on which measurements are made.** Therefore, variables are often scaled before measuring the inter-observation dissimilarities. This is particularly recommended when variables are measured in different scales (e.g: kilograms, kilometers, centimeters, ...); otherwise, the dissimilarity measures obtained will be severely affected.
- **Standardization makes the four distance measure methods — Euclidean, Manhattan, Correlation and Eisen — more similar than they would be with non-transformed data.** Note that, when the data are standardized, there is a functional relationship between the Pearson correlation coefficient & Euclidean distance.


# Strengths of k-means

- Strengths:
  - Simple: easy to understand and to implement
  - Efficient: Time complexity:  $O(tkn)$ , where  $n$  is the number of data points,  $k$  is the number of clusters, and  $t$  is the number of iterations.
  - Since both  $k$  and  $t$  are small.  $k$ -means is considered a linear algorithm.
- K-means is the most popular clustering algorithm.
- Note that: it terminates at a **local optimum** if SSE (or WCSS) is used. The **global optimum** is hard to find due to complexity.


# Weaknesses of k-means

- The algorithm is only applicable if the **mean** is defined.
  - For categorical data, *k*-mode - the centroid is represented by most frequent values.
- The user needs to specify ***k***.
- The algorithm is sensitive to **outliers**
  - Outliers are data points that are very far away from other data points.
  - Outliers could be errors in the data recording or some special data points with very different values.

# Weaknesses of k-means: Problems with outliers



(A): Undesirable clusters




(B): Ideal clusters

# Weaknesses of k-means: To deal with outliers


- One method is to remove some data points in the clustering process that are much further away from the centroids than other data points.
  - To be safe, we may want to monitor these possible outliers over a few iterations and then decide to remove them.
- Another method is to perform random sampling. Since in sampling we only choose a small subset of the data points, the chance of selecting an outlier is very small.
  - Assign the rest of the data points to the clusters by distance or similarity comparison, or classification

# Weaknesses of k-means (cont ...)


- The algorithm is sensitive to **initial seeds**.



(A). Random selection of seeds (centroids)



(B). Iteration 1




(C). Iteration 2




# Weaknesses of k-means (cont ...)

- If we use **different seeds**: good results




There are some methods to help choose good seeds

(A). Random selection of  $k$  seeds (centroids)




(B). Iteration 1




(C). Iteration 2

# Weaknesses of k-means (cont ...)

- The  $k$ -means algorithm is not suitable for discovering clusters that are not hyper-ellipsoids (or hyper-spheres).



(A): Two natural clusters



(B):  $k$ -means clusters

## Hyperparameter Tuning in K-Means Clustering



## Hyperparameters in K-Means Clustering



**Number of clusters**



**Seeds - initial values**



**Distance measures**



# How to select initial seeds?

Step 1: Select the number of clusters you want to identify in your data. This is the “K” in “K-means clustering”.

In this case, we’ll select  $K=3$ . That is to say, we want to identify 3 clusters.




There is a fancier way to select a value for “K”, but we’ll talk about that later.


Step 2: Randomly select 3 distinct data points.



Distance from the 1<sup>st</sup>  
point to the **orange**  
cluster




Step 3: Measure the distance  
between the 1<sup>st</sup> point and the three  
initial clusters.




Step 4: Assign the 1<sup>st</sup> point to the nearest cluster. In this case, the nearest cluster is the **blue** cluster.



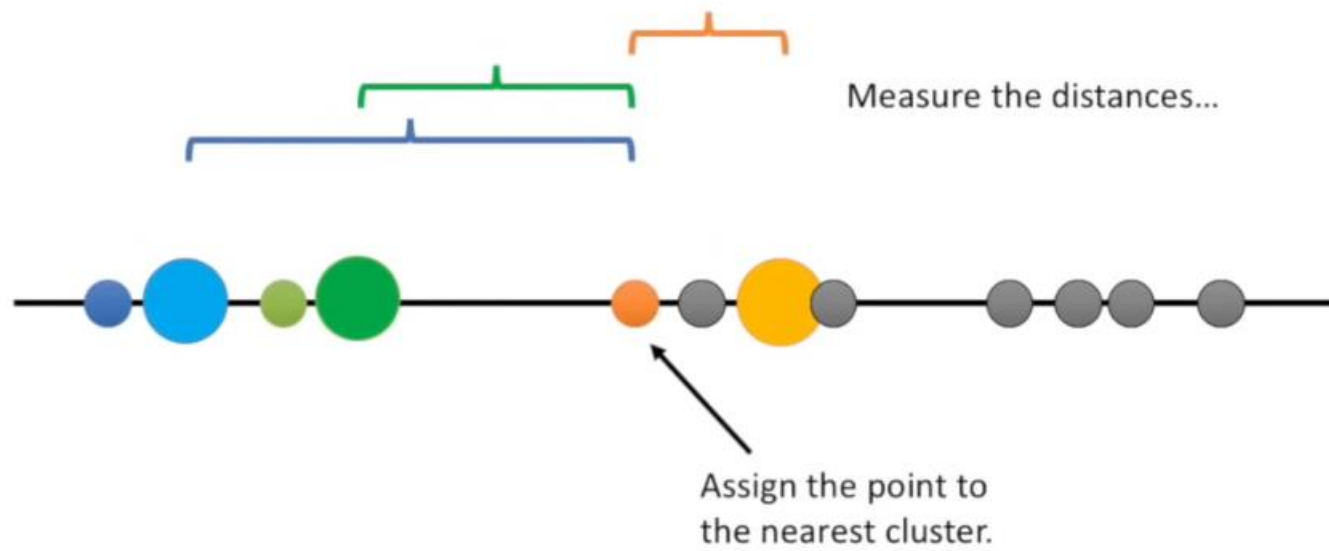
Now do the same thing for the next point.






Measure the distances...

Assign the point  
to the nearest  
cluster.






The rest of these points are  
closest to the **orange** cluster,  
so they'll go in that one, too.



The K-means clustering is pretty terrible compared to what we did by eye.



We can assess the quality of the clustering by adding up the variation within each cluster.



Since K-means clustering can't "see" the best clustering, its only option is to keep track of these clusters, and their total variance,

So, here we are again, back at the beginning.



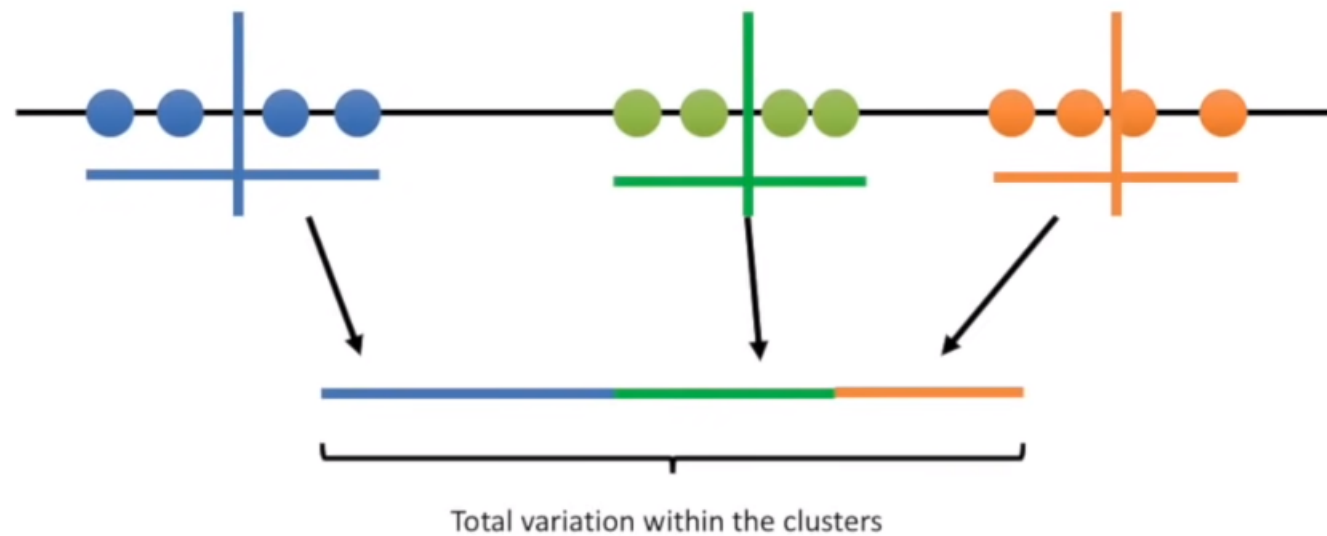
---

K-means clustering picks 3 initial clusters...






Now that the data are clustered, we sum the variation within each cluster.



At this point, K-means clustering knows that *the 2<sup>nd</sup> clustering is the best clustering so far*. But it doesn't know if it's *the best overall*, so it will do a few more clusters (it does as many as you tell it to do) and then come back and return that one if it is still the best.



1<sup>st</sup> cluster attempt:

2<sup>nd</sup> cluster attempt:

The winner!!

3<sup>rd</sup> cluster attempt:



## Number of Clusters

**K is the most important hyperparameter**

**Sometimes obvious e.g. 10 in MNIST digit classification**

**Otherwise, apply standard method to find the “best” value of K**



# Elbow Method



## Elbow Method


Pick range of candidate values of  $K$  (e.g. 1 to 10)

Calculate average distance from centroid for each value

Plot and find “elbow”




# THE ELBOW METHOD




# How to select K?

Question: How do you figure out what value to use for "K"?



# Start with $K=1$


Start with  $K = 1$



$K = 1$  is the worst case scenario. We can quantify its “badness” with the total variation.

# Now try $K=2$

Now try  $K = 2$




$K = 2$  is better, and we can quantify how much better by comparing the total variation within the 2 clusters to  $K = 1$






Now try  $K = 3$




$K = 3$  is even better! We can quantify how much better by comparing the total variation within the 3 clusters to  $K = 2$




Now try  $K = 4$



The total variation within each cluster is less than when  $K=3$




Now try  $K = 4$




The total variation within each cluster is less than when  $K=3$


Each time we add a new cluster, the total variation within each cluster is smaller than before. And when there is only one point per cluster, the variation = 0.



# THE ELBOW METHOD




# Another example



(Image by Author), Elbow Method to find optimal k

# So plot and observe!!



(Image by Author), Scatter plot of clusters formed at  $k=2, 3$ , and  $4$

# Silhouette Method



## Silhouette Method

Pick range of candidate values of  $K$  (e.g. 1 to 10)

Plot **silhouettes** for each value of  $K$

Ideal value of silhouette = 1

Worst possible value of silhouette = -1






## Silhouette Coefficient

For any point  $i$ ,  
calculate silhouette  
coefficient



## Silhouette Coefficient

Find  $a(i)$  = average  
distance of  $i$  to other  
points in **same cluster**




## Silhouette Coefficient

Find  $b(i)$  = average  
distance to **nearest  
other cluster**




## Silhouette Coefficient


Ideally,  $a(i) \ll b(i)$



## Silhouette Coefficient

If  $a(i) > b(i)$ ,  $i$  is likely misclassified






## Silhouette Coefficient

For any point  $i$

$$s(i) = \frac{b(i) - a(i)}{\text{Larger of } b(i) \text{ and } a(i)}$$

$a(i)$  = Average distance inside cluster


$b(i)$  = Average distance to nearest other cluster



Ideally  $s(i) = 1$

Ideally,  $a(i) = 0$ ,  $b(i) = \text{Infinity}$

$$s(i) = \frac{b(i) - a(i)}{\text{Larger of } b(i) \text{ and } a(i)}$$

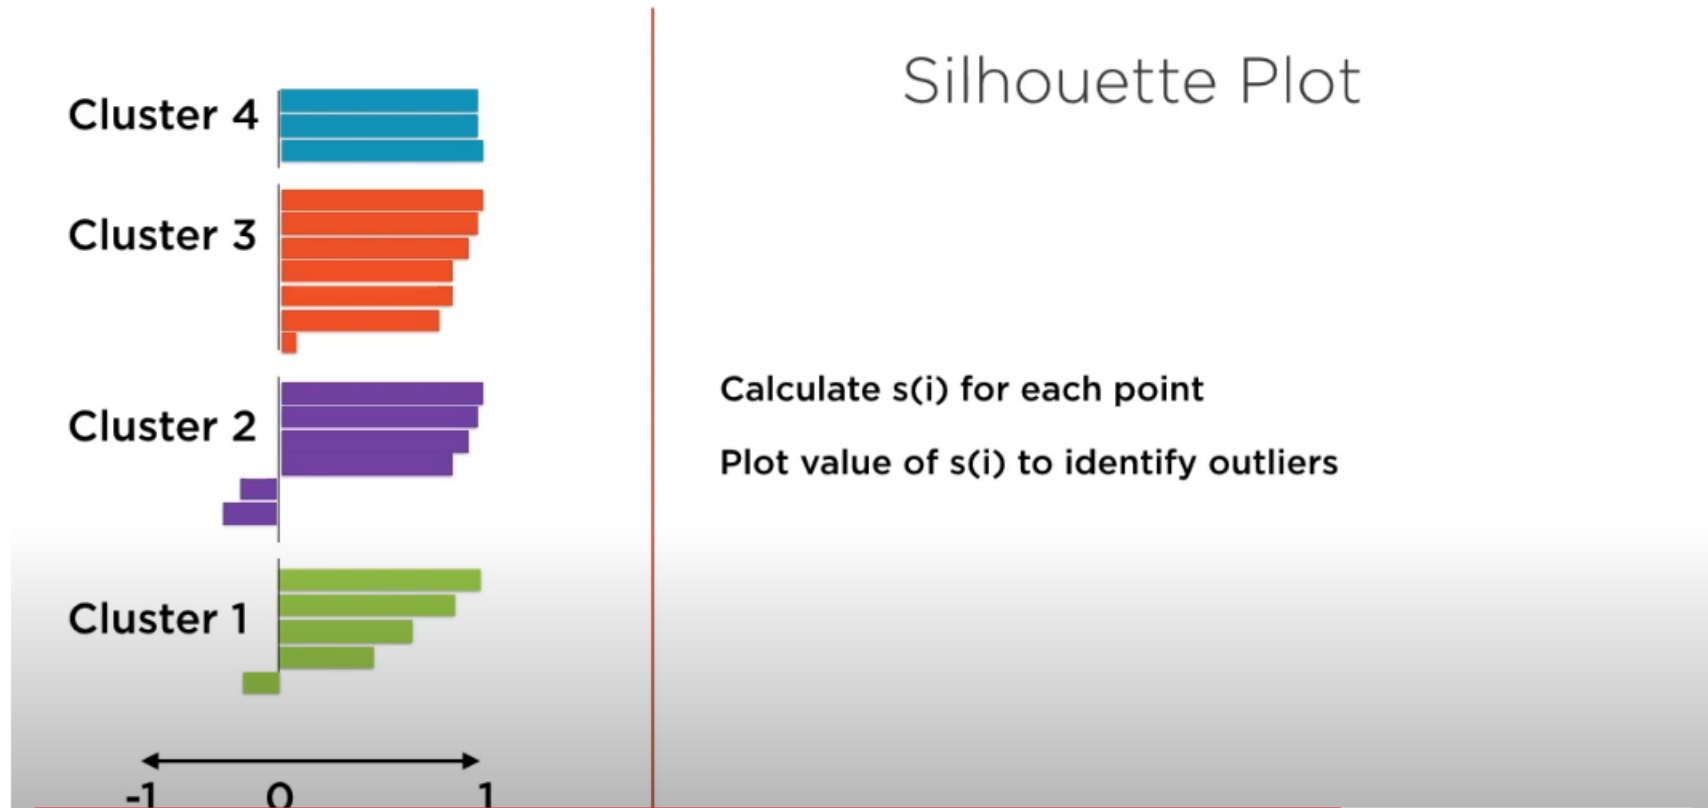



Worst-case  $s(i) = -1$

Worst case,  $a(i) = \text{Infinity}$ ,  $b(i) = 0$

$$s(i) = \frac{b(i) - a(i)}{\text{Larger of } b(i) \text{ and } a(i)} = -1$$







## Outliers

Ideally,  $s(i) = 1$

So,  $s(i) < 0$  indicates outliers




“Best”  $K$

Extend the same idea

Replicate plot for different values of  $K$

Pick  $K$  where average silhouette is closest to 1






“Best” K

Extend the same idea

Replicate plot for different values of K

Pick K where average silhouette is closest to 1



“Best”  $K$

Here  $K = 3$  is noticeably better than  $K = 4$


$K = 3$  has noticeably larger positive values




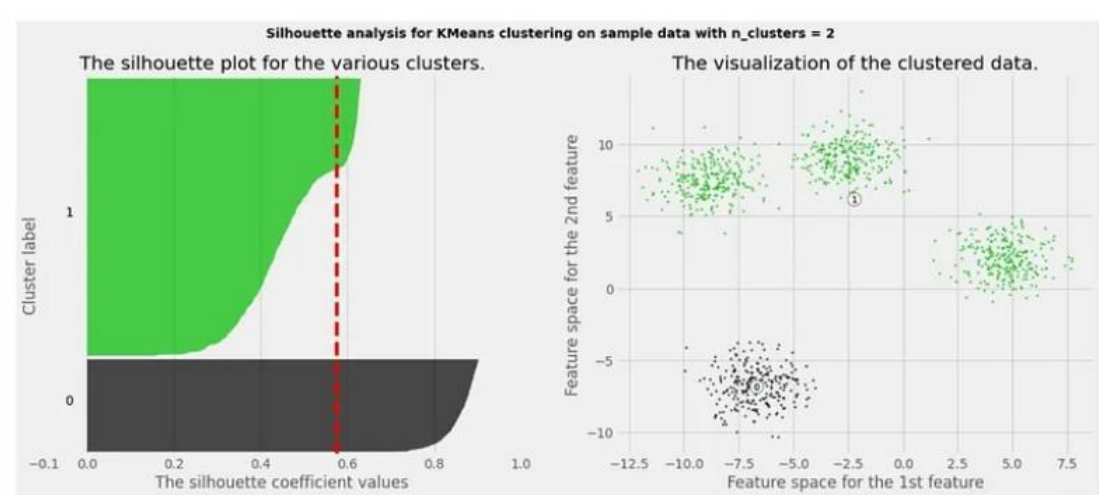
### Find the optimal value of 'k' using Silhouette Analysis:

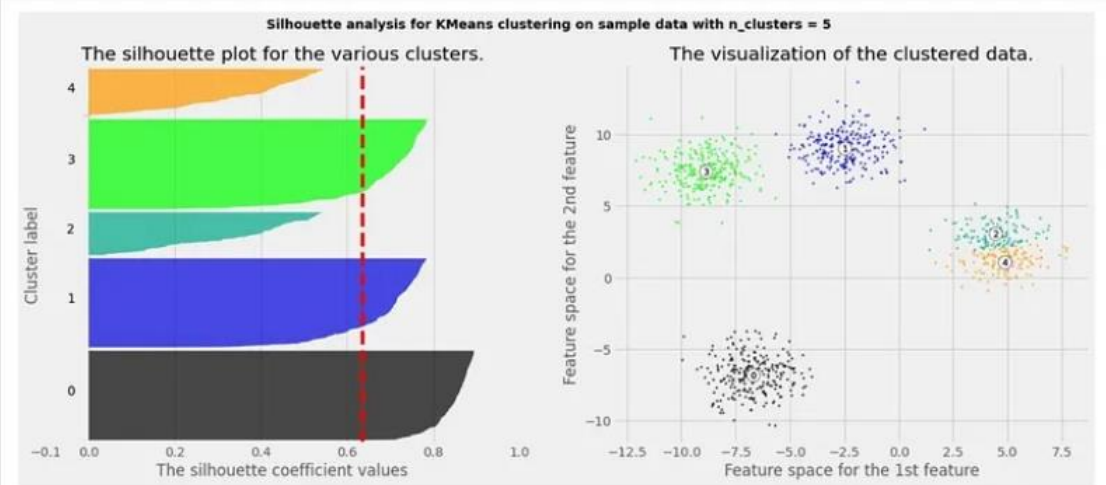
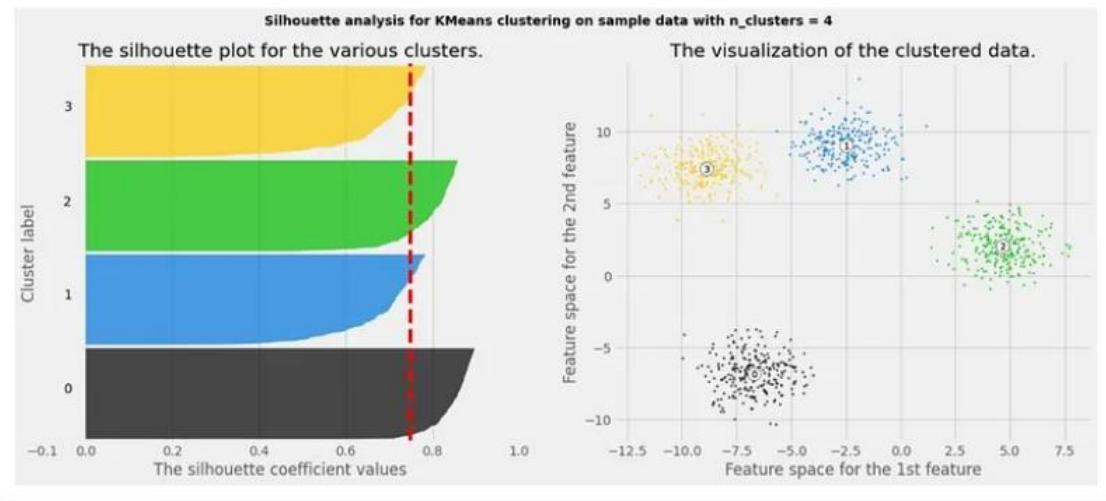
Similar to the previous Elbow method, we pick a range of candidate values of  $k$  (number of clusters), then train K-Means clustering for each of the values of  $k$ . For each k-Means clustering model represent the silhouette coefficients in a plot and observe the fluctuations and outliers of each cluster.

Top highlight



(Image by Author), Silhouette score vs the number of clusters







# K-means summary

- Despite weaknesses, *k*-means is still the most popular algorithm due to its simplicity, efficiency and
  - other clustering algorithms have their own lists of weaknesses (C-Means, Dbscan,...)
- No clear evidence that any other clustering algorithm performs better in general
  - although they may be more suitable for some specific types of data or applications.
- Comparing different clustering algorithms is a difficult task. No one knows the correct clusters!